

Introduction to Map Reduce

Continuum Analytics



Large Relative to What?

- Can't fit into Excel

Large Relative to What?

- Can't fit into Excel
 - Increase Memory

Large Relative to What?

- Can't fit into Excel
 - Increase Memory
- Can't fit into R

Large Relative to What?

- Can't fit into Excel
 - Increase Memory
- Can't fit into R
 - Increase Memory

Large Relative to What?

- Can't fit into Excel
 - Increase Memory
- Can't fit into R
 - Increase Memory
- Can't fit into Memory
 - Increase Memory

Large Relative to What?

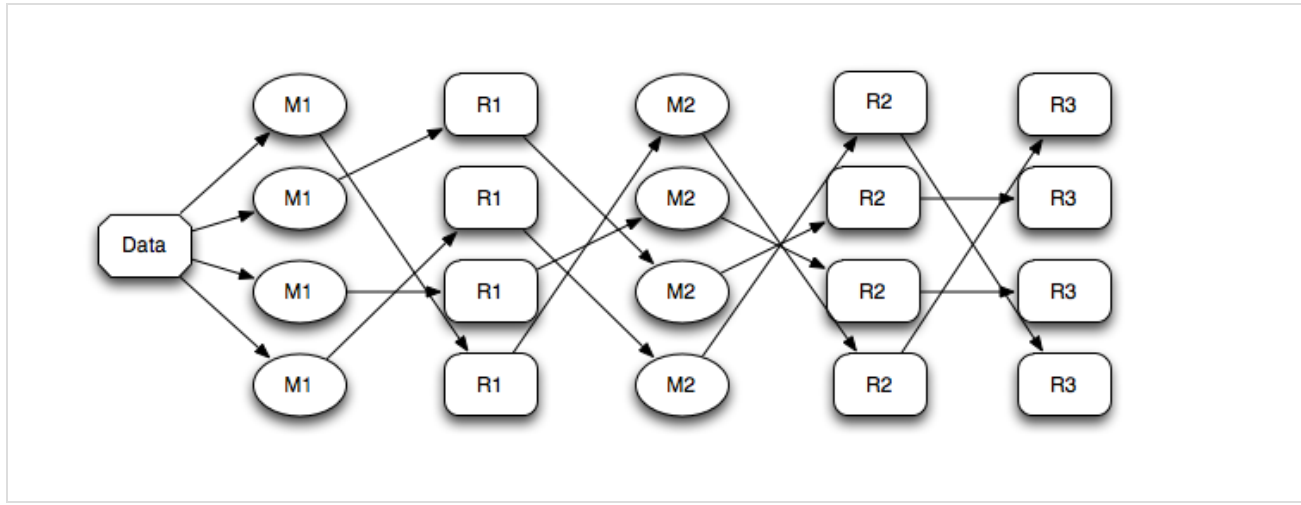
- Can't fit into Excel
 - Increase Memory
- Can't fit into R
 - Increase Memory
- Can't fit into Memory
 - Increase Memory
- Can't fit on a single disk
 - Distributed Filesystem: SAN, HDFS/DDFS, AWS: S3, Redshift, etc.

MapReduce

Framework to help solve the problem of distributed computation for distributed data

- A mass of data: records
- Split/**Map** records into key-values pairs
- Collect/Partition kv pairs (Optional Sort)
- Buckets are passed to **Reduce** function
- Result is returned

MapReduce Workflow



- Push Code to Data
- Lots of Network Traffic

MR Implementations

- Disco: Python + Erlang
 - Distributed FileSystem: DDFS
- Hadoop: Java
 - Streaming with Python
 - Dumbo
 - MRJob
 - Hadoopy

MapReduce: It's a Party



Buddies Included


- NumPy
- SciPy
- pandas
- scikits-learn
- OpenCV
- ...

Canonical Example

```
1 from disco.job import Job
2 from disco.core import result_iterator
3
4 class WordCount(Job):
5
6     partitions = 3
7     input=["sherlock.txt","poiroit.txt","clouseau.txt"]
8
9     @staticmethod
10    def map(line, params):
11        import string
12        for word in line.split():
13            yield word, 1
14
15    @staticmethod
16    def reduce(iter, params):
17        from disco.util import kvgroup
18        for word, counts in kvgroup(sorted(iter)):
19            yield word, sum(counts)
20
21 if __name__ == "__main__":
22     from disco_words import WordCount
23
24     wordcount = WordCount().run()
25
26     for (word, counts) in result_iterator(wordcount.wait(show=True)):
27         print word, counts
28
```


Demo 1


Chaining Jobs



MapReduce Thoughts

- Data Cleansing
 - Everyone's pain point
- Task Deconstruction
 - Good for code management
 - Hides -- in a good way -- data management
- Can Be Inefficient
 - Network traffic
 - Job organization






**83% of the radiologists
shown CT missed the
gorilla**

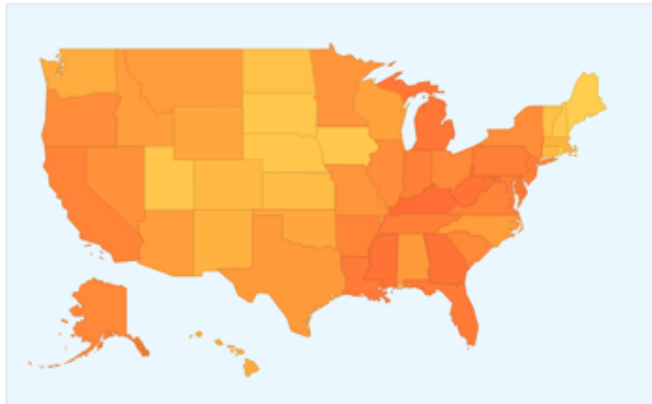
Google Flu

Explore flu trends - United States

We've found that certain search terms are good indicators of flu activity. Google Flu Trends uses aggregated Google search data to estimate flu activity. [Learn more >](#)



States | [Cities](#) (Experimental)




- Data Mining
- Faster than CDC

Google Get's It Wrong

FEVER PEAKS

A comparison of three different methods of measuring the proportion of the US population with an influenza-like illness.



- Typically, prediction is great!
- This year not so much
- Google: No comment!
- Feedback mechanism from hype-up media

Data Philosophy

- Invisible Gorillas will stay Invisible
 - Inattentional Blindness
- Machine Learning without Oversight
 - Turnkey analytics is dangerous
- Good Analysis
 - Requires iterative exploration
 - Peer review and collaboration

Canonical Example

```
class WordCount(Job):
    partitions = 3
    input=["sherlock.txt", "poirot.txt", "clouseau.txt"]

    @staticmethod
    def map(line, params):
        import string
        for word in line.split():
            yield word, 1

    @staticmethod
    def reduce(iter, params):
        from disco.util import kvgroup
        for word, counts in kvgroup(sorted(iter)):
            yield word, sum(counts)

if __name__ == "__main__":
    from count_words import WordCount

    for (word, counts) in result_iterator(WordCount.wait(show=True)):
        print word, counts
```